

Solving high-dimensional Hamilton-Jacobi-Bellman Equations for Optimal Feedback Control via Adaptive Deep Learning Approach

Elisa Giesecke

Supervisors: Prof. Dr. Falk Hante, Prof. Dr. Daniel Walter

Humboldt-Universität zu Berlin
Mathematisch-Naturwissenschaftliche Fakultät
Institut für Mathematik

February 22, 2024

- 1 Introduction
- 2 Optimal Control Theory
- 3 Adaptive Deep Learning
- 4 Application to Reaction-Diffusion System
- 5 Numerical Experiments
- 6 Conclusion and Future Research

- 1 Introduction
- 2 Optimal Control Theory
- 3 Adaptive Deep Learning
- 4 Application to Reaction-Diffusion System
- 5 Numerical Experiments
- 6 Conclusion and Future Research

Goal: designing feedback controllers for OCPs with non-linear ODEs
→ solving Hamilton-Jacobi-Bellman equations

Goal: designing feedback controllers for OCPs with non-linear ODEs
→ solving Hamilton-Jacobi-Bellman equations

"curse of dimensionality"

Motivation

Goal: designing feedback controllers for OCPs with non-linear ODEs
→ solving Hamilton-Jacobi-Bellman equations

"curse of dimensionality"

limited
scalability
w. r. t. space
dimension

restrictive
structure of
dynamical
system

solution
validity in
local neigh-
borhood

difficult
verification
of solution
accuracy

Motivation

Goal: designing feedback controllers for OCPs with non-linear ODEs
→ solving Hamilton-Jacobi-Bellman equations

"curse of dimensionality"

limited
scalability
w. r. t. space
dimension

restrictive
structure of
dynamical
system

solution
validity in
local neigh-
borhood

difficult
verification
of solution
accuracy

→ developing data-driven approaches based on neural networks

Advancing adaptive deep learning approach from
[Nakamura-Zimmerer et al., 2021]:

Advancing adaptive deep learning approach from
[Nakamura-Zimmerer et al., 2021]:

- value function approximator with exact final condition

Advancing adaptive deep learning approach from
[Nakamura-Zimmerer et al., 2021]:

- value function approximator with exact final condition
- residual network architecture

Advancing adaptive deep learning approach from [Nakamura-Zimmerer et al., 2021]:

- value function approximator with exact final condition
- residual network architecture
- first-order stochastic optimizer Adam

Advancing adaptive deep learning approach from [Nakamura-Zimmerer et al., 2021]:

- value function approximator with exact final condition
- residual network architecture
- first-order stochastic optimizer Adam
- physics-informed gradient regularization and L^2 regularization

Advancing adaptive deep learning approach from [Nakamura-Zimmerer et al., 2021]:

- value function approximator with exact final condition
- residual network architecture
- first-order stochastic optimizer Adam
- physics-informed gradient regularization and L^2 regularization
- early stopping technique combined with variance estimation based sample size selection

Advancing adaptive deep learning approach from [Nakamura-Zimmerer et al., 2021]:

- value function approximator with exact final condition
- residual network architecture
- first-order stochastic optimizer Adam
- physics-informed gradient regularization and L^2 regularization
- early stopping technique combined with variance estimation based sample size selection

Implementing algorithm applied to controlled reaction-diffusion system for numerical analysis:

<https://github.com/ElisaGiesecke/AdaDL-for-HJB-Equations>.

- 1 Introduction
- 2 Optimal Control Theory
- 3 Adaptive Deep Learning
- 4 Application to Reaction-Diffusion System
- 5 Numerical Experiments
- 6 Conclusion and Future Research

Problem Formulation

- initial time $t_0 \in \mathbb{R}$, final time $t_f \in \mathbb{R}$
- control $\mathbf{u} : [t_0, t_f] \rightarrow \mathbb{R}^m$, $\mathcal{U} := L^\infty(t_0, t_f; \mathbb{R}^m)$
- state $\mathbf{x} : [t_0, t_f] \rightarrow \mathbb{R}^n$, $\mathcal{X} := W^{1,\infty}(t_0, t_f; \mathbb{R}^n)$

Problem (fixed-time free-endpoint optimal control problem)

Problem Formulation

- initial time $t_0 \in \mathbb{R}$, final time $t_f \in \mathbb{R}$
- control $\mathbf{u} : [t_0, t_f] \rightarrow \mathbb{R}^m$, $\mathcal{U} := L^\infty(t_0, t_f; \mathbb{R}^m)$
- state $\mathbf{x} : [t_0, t_f] \rightarrow \mathbb{R}^n$, $\mathcal{X} := W^{1,\infty}(t_0, t_f; \mathbb{R}^n)$
- dynamics $f : [t_0, t_f] \times \mathbb{R}^m \times \mathbb{R}^n \rightarrow \mathbb{R}^n$
- initial condition $x_0 \in \mathbb{R}^n$

Problem (fixed-time free-endpoint optimal control problem)

$$\dot{\mathbf{x}}(t) = f(t, \mathbf{u}(t), \mathbf{x}(t)) \quad \text{for a. a. } t \in [t_0, t_f]; \quad \mathbf{x}(0) = x_0;$$

Problem Formulation

- initial time $t_0 \in \mathbb{R}$, final time $t_f \in \mathbb{R}$
- control $\mathbf{u} : [t_0, t_f] \rightarrow \mathbb{R}^m$, $\mathcal{U} := L^\infty(t_0, t_f; \mathbb{R}^m)$
- state $\mathbf{x} : [t_0, t_f] \rightarrow \mathbb{R}^n$, $\mathcal{X} := W^{1,\infty}(t_0, t_f; \mathbb{R}^n)$
- dynamics $f : [t_0, t_f] \times \mathbb{R}^m \times \mathbb{R}^n \rightarrow \mathbb{R}^n$
- initial condition $x_0 \in \mathbb{R}^n$
- cost function $J : \mathcal{U} \times \mathcal{X} \rightarrow \mathbb{R}$
- running cost $\psi : [t_0, t_f] \times \mathbb{R}^m \times \mathbb{R}^n \rightarrow \mathbb{R}$, final cost $\phi : \mathbb{R}^n \rightarrow \mathbb{R}$

Problem (fixed-time free-endpoint optimal control problem)

$$\min_{(\mathbf{u}, \mathbf{x}) \in \mathcal{U} \times \mathcal{X}} J(\mathbf{u}, \mathbf{x}) := \int_{t_0}^{t_f} \psi(t, \mathbf{u}(t), \mathbf{x}(t)) \, dt + \phi(\mathbf{x}(t_f))$$

subject to $\dot{\mathbf{x}}(t) = f(t, \mathbf{u}(t), \mathbf{x}(t))$ for a. a. $t \in [t_0, t_f]$; $\mathbf{x}(0) = x_0$;

Problem Formulation

- initial time $t_0 \in \mathbb{R}$, final time $t_f \in \mathbb{R}$
- control $\mathbf{u} : [t_0, t_f] \rightarrow \mathbb{R}^m$, $\mathcal{U} := L^\infty(t_0, t_f; \mathbb{R}^m)$
- state $\mathbf{x} : [t_0, t_f] \rightarrow \mathbb{R}^n$, $\mathcal{X} := W^{1,\infty}(t_0, t_f; \mathbb{R}^n)$
- dynamics $f : [t_0, t_f] \times \mathbb{R}^m \times \mathbb{R}^n \rightarrow \mathbb{R}^n$
- initial condition $x_0 \in \mathbb{R}^n$
- cost function $J : \mathcal{U} \times \mathcal{X} \rightarrow \mathbb{R}$
- running cost $\psi : [t_0, t_f] \times \mathbb{R}^m \times \mathbb{R}^n \rightarrow \mathbb{R}$, final cost $\phi : \mathbb{R}^n \rightarrow \mathbb{R}$
- control set $U \subseteq \mathbb{R}^m$

Problem (fixed-time free-endpoint optimal control problem)

$$\min_{(\mathbf{u}, \mathbf{x}) \in \mathcal{U} \times \mathcal{X}} J(\mathbf{u}, \mathbf{x}) := \int_{t_0}^{t_f} \psi(t, \mathbf{u}(t), \mathbf{x}(t)) \, dt + \phi(\mathbf{x}(t_f))$$

subject to $\dot{\mathbf{x}}(t) = f(t, \mathbf{u}(t), \mathbf{x}(t))$ for a. a. $t \in [t_0, t_f]$; $\mathbf{x}(0) = x_0$;
 $\mathbf{u}(t) \in U$ for a. a. $t \in [t_0, t_f]$.

Derivation of Conditions for Optimality

variational approach

necessary conditions via
Pontryagin's Minimum Principle
based on adjoint state as
Lagrange multiplier
→ open-loop control

Derivation of Conditions for Optimality

variational approach

necessary conditions via
Pontryagin's Minimum Principle
based on adjoint state as
Lagrange multiplier
→ open-loop control

dynamic programming approach

sufficient conditions based on
value function as solution of
Hamilton-Jacobi-Bellman
equation
→ closed-loop feedback control

Derivation of Conditions for Optimality

variational approach

necessary conditions via
Pontryagin's Minimum Principle
based on adjoint state as
Lagrange multiplier
→ open-loop control

dynamic programming approach

sufficient conditions based on
value function as solution of
Hamilton-Jacobi-Bellman
equation
→ closed-loop feedback control

two-point boundary value problem for supervised learning

candidate open-loop solutions for data generation
value function learned by neural network

Open-loop Control via Pontryagin's Minimum Principle

Definition (Hamiltonian function)

$$H : \mathbb{R} \times \mathbb{R}^m \times \mathbb{R}^n \times \mathbb{R}^n \rightarrow \mathbb{R}, \quad H(t, u, x, p) := \psi(t, u, x) + p \cdot f(t, u, x)$$

Open-loop Control via Pontryagin's Minimum Principle

Definition (Hamiltonian function)

$$H : \mathbb{R} \times \mathbb{R}^m \times \mathbb{R}^n \times \mathbb{R}^n \rightarrow \mathbb{R}, \quad H(t, u, x, p) := \psi(t, u, x) + p \cdot f(t, u, x)$$

Theorem (Pontryagin's minimum principle) [Schättler et al., 2012]

Let $\mathbf{u}^* \in \mathcal{U}$ be a globally optimal control of the OCP with initial condition x_0 . Furthermore, let $\mathbf{x}^* \in \mathcal{X}$ be the corresponding optimal state trajectory and $\mathbf{p}^* \in \mathcal{X}$ the associated costate, i. e., the solutions to the **Hamiltonian system**

$$\begin{cases} \dot{\mathbf{x}}(t) = H_p(t, \mathbf{u}^*(t; x_0), \mathbf{x}(t), \mathbf{p}(t)) & \text{for a. a. } t; \quad \mathbf{x}(t_0) = x_0, \\ \dot{\mathbf{p}}(t) = -H_x(t, \mathbf{u}^*(t; x_0), \mathbf{x}(t), \mathbf{p}(t)) & \text{for a. a. } t; \quad \mathbf{p}(t_f) = \phi_x(\mathbf{x}(t_f)). \end{cases}$$

Open-loop Control via Pontryagin's Minimum Principle

Definition (Hamiltonian function)

$$H : \mathbb{R} \times \mathbb{R}^m \times \mathbb{R}^n \times \mathbb{R}^n \rightarrow \mathbb{R}, \quad H(t, u, x, p) := \psi(t, u, x) + p \cdot f(t, u, x)$$

Theorem (Pontryagin's minimum principle) [Schättler et al., 2012]

Let $\mathbf{u}^* \in \mathcal{U}$ be a globally optimal control of the OCP with initial condition x_0 . Furthermore, let $\mathbf{x}^* \in \mathcal{X}$ be the corresponding optimal state trajectory and $\mathbf{p}^* \in \mathcal{X}$ the associated costate, i. e., the solutions to the Hamiltonian system

$$\begin{cases} \dot{\mathbf{x}}(t) = H_p(t, \mathbf{u}^*(t; x_0), \mathbf{x}(t), \mathbf{p}(t)) & \text{for a. a. } t; \quad \mathbf{x}(t_0) = x_0, \\ \dot{\mathbf{p}}(t) = -H_x(t, \mathbf{u}^*(t; x_0), \mathbf{x}(t), \mathbf{p}(t)) & \text{for a. a. } t; \quad \mathbf{p}(t_f) = \phi_x(\mathbf{x}(t_f)). \end{cases}$$

Then the **Hamiltonian minimization condition** below holds:

$$\mathbf{u}^*(t) \in \underset{u \in U}{\operatorname{argmin}} H(t, u, \mathbf{x}^*(t), \mathbf{p}^*(t)) \quad \text{for a. a. } t.$$

Theorem [Liberzon, 2012]

Consider the OCP and suppose that a continuously differentiable function $V : [t_0, t_f] \times \mathbb{R}^n \rightarrow \mathbb{R}$ is a classical solution to the **HJB equation** with its final condition, i. e., satisfying

$$\begin{cases} -V_t(t, x) = \inf_{u \in U} H(t, u, x, V_x(t, x)) & \text{for all } t \text{ and } x, \\ V(t_f, x) = \phi(x) & \text{for all } x. \end{cases}$$

Feedback Control via Hamilton-Jacobi-Bellman Equation

Theorem [Liberzon, 2012]

Consider the OCP and suppose that a continuously differentiable function $V : [t_0, t_f] \times \mathbb{R}^n \rightarrow \mathbb{R}$ is a classical solution to the HJB equation with its final condition, i. e., satisfying

$$\begin{cases} -V_t(t, x) = \inf_{u \in U} H(t, u, x, V_x(t, x)) & \text{for all } t \text{ and } x, \\ V(t_f, x) = \phi(x) & \text{for all } x. \end{cases}$$

Furthermore, suppose that there exists a control $\mathbf{u}^* \in \mathcal{U}$ and a state $\mathbf{x}^* \in \mathcal{X}$ fulfilling the state equation with initial condition, for which the **Hamiltonian minimization condition**

$$\mathbf{u}^*(t) \in \underset{u \in U}{\operatorname{argmin}} H(t, u, \mathbf{x}^*(t), V_x(t, \mathbf{x}^*(t))) \quad \text{for all } t$$

holds. Then \mathbf{u}^* is a globally optimal control.

Relation between PMP and HJB approach

- costate as gradient of the value function

$$\mathbf{p}^*(t) = V_x(t, \mathbf{x}^*(t)), \quad \text{assuming } V \in \mathcal{C}^2$$

Relation between PMP and HJB approach

- costate as gradient of the value function

$$\mathbf{p}^*(t) = V_x(t, \mathbf{x}^*(t)), \quad \text{assuming } V \in \mathcal{C}^2$$

- Hamiltonian system as characteristics of the HJB equation

$$\begin{cases} \dot{\mathbf{x}}(t) = H_p(t, \mathbf{u}^*(t; x_0), \mathbf{x}(t), \mathbf{p}(t)) & \text{for a. a. } t; \quad \mathbf{x}(t_0) = x_0, \\ \dot{\mathbf{p}}(t) = -H_x(t, \mathbf{u}^*(t; x_0), \mathbf{x}(t), \mathbf{p}(t)) & \text{for a. a. } t; \quad \mathbf{p}(t_f) = \phi_x(\mathbf{x}(t_f)), \end{cases}$$

Relation between PMP and HJB approach

- costate as gradient of the value function

$$\mathbf{p}^*(t) = V_x(t, \mathbf{x}^*(t)), \quad \text{assuming } V \in \mathcal{C}^2$$

- Hamiltonian system as characteristics of the HJB equation

$$\begin{cases} \dot{\mathbf{x}}(t) = H_p(t, \mathbf{u}^*(t; x_0), \mathbf{x}(t), \mathbf{p}(t)) & \text{for a. a. } t; \quad \mathbf{x}(t_0) = x_0, \\ \dot{\mathbf{p}}(t) = -H_x(t, \mathbf{u}^*(t; x_0), \mathbf{x}(t), \mathbf{p}(t)) & \text{for a. a. } t; \quad \mathbf{p}(t_f) = \phi_x(\mathbf{x}(t_f)), \\ \dot{\mathbf{v}}(t) = -\psi(t, \mathbf{u}^*(t; x_0), \mathbf{x}(t)) & \text{for a. a. } t; \quad \mathbf{v}(t_f) = \phi(\mathbf{x}(t_f)), \end{cases}$$

extended by ODE describing the evolution of the value function along trajectory, denoting $\mathbf{v}(t) := V(t, \mathbf{x}^*(t))$, see [Kang and Wilcox, 2015]

- 1 Introduction
- 2 Optimal Control Theory
- 3 Adaptive Deep Learning**
- 4 Application to Reaction-Diffusion System
- 5 Numerical Experiments
- 6 Conclusion and Future Research

Value Function Approximation for NN Feedback Control

- value function approximator V^{NN} with neural network function F^{NN}
→ guaranteeing satisfaction of final condition

$$V^{\text{NN}}(t, x) := F^{\text{NN}}(t, x) - F^{\text{NN}}(t_f, x) + \phi(x) \quad \text{for all } t \text{ and } x$$

Value Function Approximation for NN Feedback Control

- value function approximator V^{NN} with neural network function F^{NN}
→ guaranteeing satisfaction of final condition

$$V^{\text{NN}}(t, x) := F^{\text{NN}}(t, x) - F^{\text{NN}}(t_f, x) + \phi(x) \quad \text{for all } t \text{ and } x$$

- NN feedback controller \mathbf{u}^{NN}
→ exploiting Hamiltonian minimization condition

$$\mathbf{u}^{\text{NN}}(t, x) \in \underset{u \in \mathbb{R}^m}{\operatorname{argmin}} H(t, u, x, V_x^{\text{NN}}(t, x)) \quad \text{for all } t \text{ and } x$$

Residual Network Architecture

$$F^{\text{NN}} = f^{[L]} \circ \dots \circ f^{[0]}, \quad \theta = \{W^{[l]}, b^{[l]}\}_{l=0}^L$$


- residual layers

$$f^{[l]}(y) = y + \sigma(W^{[l]}y + b^{[l]})$$

- weights $W^{[l]} \in \mathbb{R}^{d_{l+1} \times d_l}$ and biases $b^{[l]} \in \mathbb{R}^{d_{l+1}}$

- softplus activation function


$$\sigma(x) = \log(1 + \exp(x))$$



Residual Network Architecture

$$F^{\text{NN}} = f^{[L]} \circ \dots \circ f^{[0]}, \quad \theta = \{W^{[l]}, b^{[l]}\}_{l=0}^L$$

- residual layers
 $f^{[l]}(y) = y + \sigma(W^{[l]}y + b^{[l]})$
- weights $W^{[l]} \in \mathbb{R}^{d_{l+1} \times d_l}$ and biases $b^{[l]} \in \mathbb{R}^{d_{l+1}}$
- softplus activation function
 $\sigma(x) = \log(1 + \exp(x))$



- avoiding vanishing gradients and representational bottlenecks
- interpreting as explicit forward Euler discretization of neural ODE

Physics-Informed Training with Adam

- training data $D_{\text{train}} = \{(t^{(k)}, x^{(k)}), V^{(k)}\}_{k=1}^K$

Problem (physics-informed deep learning problem) [Raissi et al., 2019]

Physics-Informed Training with Adam

- training data $D_{\text{train}} = \{(t^{(k)}, x^{(k)}), V^{(k)}\}_{k=1}^K$

- mean squared error loss

$$\text{MSE}_V(\theta; D_{\text{train}}) := \frac{1}{K} \sum_{k=1}^K |V^{\text{NN}}(t^{(k)}, x^{(k)}; \theta) - V^{(k)}|^2$$

Problem (physics-informed deep learning problem) [Raissi et al., 2019]

$$\min_{\theta} \text{Cost}(\theta; D_{\text{train}}) := \text{MSE}_V(\theta; D_{\text{train}})$$

Physics-Informed Training with Adam

- training data $D_{\text{train}} = \{(t^{(k)}, x^{(k)}), (V^{(k)}, p^{(k)})\}_{k=1}^K$
- mean squared error loss
$$\text{MSE}_V(\theta; D_{\text{train}}) := \frac{1}{K} \sum_{k=1}^K |V^{\text{NN}}(t^{(k)}, x^{(k)}; \theta) - V^{(k)}|^2$$
- gradient regularization
$$\text{MSE}_p(\theta; D_{\text{train}}) := \frac{1}{K} \sum_{k=1}^K \|V_x^{\text{NN}}(t^{(k)}, x^{(k)}; \theta) - p^{(k)}\|_2^2, \quad \lambda > 0$$

Problem (physics-informed deep learning problem) [Raissi et al., 2019]

$$\min_{\theta} \text{Cost}(\theta; D_{\text{train}}) := \text{MSE}_V(\theta; D_{\text{train}}) + \lambda \cdot \text{MSE}_p(\theta; D_{\text{train}})$$

Physics-Informed Training with Adam

- training data $D_{\text{train}} = \{(t^{(k)}, x^{(k)}), (V^{(k)}, p^{(k)})\}_{k=1}^K$
- mean squared error loss
$$\text{MSE}_V(\theta; D_{\text{train}}) := \frac{1}{K} \sum_{k=1}^K |V^{\text{NN}}(t^{(k)}, x^{(k)}; \theta) - V^{(k)}|^2$$
- gradient regularization
$$\text{MSE}_p(\theta; D_{\text{train}}) := \frac{1}{K} \sum_{k=1}^K \|V_x^{\text{NN}}(t^{(k)}, x^{(k)}; \theta) - p^{(k)}\|_2^2, \quad \lambda > 0$$

Problem (physics-informed deep learning problem) [Raissi et al., 2019]

$$\min_{\theta} \text{Cost}(\theta; D_{\text{train}}) := \text{MSE}_V(\theta; D_{\text{train}}) + \lambda \cdot \text{MSE}_p(\theta; D_{\text{train}})$$

- complementing gradient regularization with weight decay
- optimizing via batch gradient descent with adaptive learning rates

Progressive and Adaptive Data Generation

Employing spatially causality-free solver for BVP

$$\begin{cases} \dot{\mathbf{x}}(t) = H_p(t, \mathbf{u}^*(t; x_0), \mathbf{x}(t), \mathbf{p}(t)) & \text{for a. a. } t; \quad \mathbf{x}(t_0) = x_0, \\ \dot{\mathbf{p}}(t) = -H_x(t, \mathbf{u}^*(t; x_0), \mathbf{x}(t), \mathbf{p}(t)) & \text{for a. a. } t; \quad \mathbf{p}(t_f) = \phi_x(\mathbf{x}(t_f)), \\ \dot{\mathbf{v}}(t) = -\psi(t, \mathbf{u}^*(t; x_0), \mathbf{x}(t)) & \text{for a. a. } t; \quad \mathbf{v}(t_f) = \phi(\mathbf{x}(t_f)), \end{cases}$$

for generating training, validation and test data sets

Employing spatially causality-free solver for BVP

$$\begin{cases} \dot{\mathbf{x}}(t) = H_p(t, \mathbf{u}^*(t; x_0), \mathbf{x}(t), \mathbf{p}(t)) & \text{for a. a. } t; \quad \mathbf{x}(t_0) = x_0, \\ \dot{\mathbf{p}}(t) = -H_x(t, \mathbf{u}^*(t; x_0), \mathbf{x}(t), \mathbf{p}(t)) & \text{for a. a. } t; \quad \mathbf{p}(t_f) = \phi_x(\mathbf{x}(t_f)), \\ \dot{\mathbf{v}}(t) = -\psi(t, \mathbf{u}^*(t; x_0), \mathbf{x}(t)) & \text{for a. a. } t; \quad \mathbf{v}(t_f) = \phi(\mathbf{x}(t_f)), \end{cases}$$

for generating training, validation and test data sets

- ability to generate data in selected regions
→ uniform or adaptive sampling

Employing spatially causality-free solver for BVP

$$\begin{cases} \dot{\mathbf{x}}(t) = H_p(t, \mathbf{u}^*(t; x_0), \mathbf{x}(t), \mathbf{p}(t)) & \text{for a. a. } t; \quad \mathbf{x}(t_0) = x_0, \\ \dot{\mathbf{p}}(t) = -H_x(t, \mathbf{u}^*(t; x_0), \mathbf{x}(t), \mathbf{p}(t)) & \text{for a. a. } t; \quad \mathbf{p}(t_f) = \phi_x(\mathbf{x}(t_f)), \\ \dot{\mathbf{v}}(t) = -\psi(t, \mathbf{u}^*(t; x_0), \mathbf{x}(t)) & \text{for a. a. } t; \quad \mathbf{v}(t_f) = \phi(\mathbf{x}(t_f)), \end{cases}$$

for generating training, validation and test data sets

- ability to generate data in selected regions
→ uniform or adaptive sampling
- sensitivity to initialization of \mathbf{x} , \mathbf{p} and \mathbf{v}
→ time-marching or NN warm start

Initialization and Sampling Techniques

Before training: initial training set and validation set


D_{train}^0

D_{val}

- uniform sampling: drawing initial conditions uniformly at random
- time-marching: solving over time intervals of increasing length

Initialization and Sampling Techniques


During training: **adaptively augmented training sets**



- adaptive sampling: selecting initial conditions where predicted value function gradient is large
- NN warm start: simulating with partially trained neural network

Initialization and Sampling Techniques

After training: **test set**



- sampling as desired: drawing initial conditions randomly or in region of interest
- NN warm start: simulating with fully trained neural network

Empirical Validation of Model Accuracy

Error metrics: for some data set $D = \{(t^{(k)}, x^{(k)}), (V^{(k)}, p^{(k)})\}_{k=1}^K$

- $\text{RMAE}_V(\theta; D) := \frac{\sum_{k=1}^K |V^{\text{NN}}(t^{(k)}, x^{(k)}; \theta) - V^{(k)}|}{\sum_{k=1}^K |V^{(k)}|}$

Empirical Validation of Model Accuracy

Error metrics: for some data set $D = \{(t^{(k)}, x^{(k)}), (V^{(k)}, p^{(k)})\}_{k=1}^K$

- $\text{RMAE}_V(\theta; D) := \frac{\sum_{k=1}^K |V^{\text{NN}}(t^{(k)}, x^{(k)}; \theta) - V^{(k)}|}{\sum_{k=1}^K |V^{(k)}|}$
- $\text{RML}^1_p(\theta; D) := \frac{\sum_{k=1}^K \|V_x^{\text{NN}}(t^{(k)}, x^{(k)}; \theta) - p^{(k)}\|_1}{\sum_{k=1}^K \|p^{(k)}\|_1}$

Empirical Validation of Model Accuracy

Error metrics: for some data set $D = \{(t^{(k)}, x^{(k)}), (V^{(k)}, p^{(k)})\}_{k=1}^K$

- $\text{RMAE}_V(\theta; D) := \frac{\sum_{k=1}^K |V^{\text{NN}}(t^{(k)}, x^{(k)}; \theta) - V^{(k)}|}{\sum_{k=1}^K |V^{(k)}|}$
- $\text{RML}^1_p(\theta; D) := \frac{\sum_{k=1}^K \|V_x^{\text{NN}}(t^{(k)}, x^{(k)}; \theta) - p^{(k)}\|_1}{\sum_{k=1}^K \|p^{(k)}\|_1}$
- $\text{RMAE}_u(\theta; D) := \frac{\sum_{k=1}^K |u^{\text{NN}}(t^{(k)}, x^{(k)}; \theta) - u^{(k)}|}{\sum_{k=1}^K |u^{(k)}|}, \quad u^{(k)} = \mathbf{u}(t^{(k)})$

Empirical Validation of Model Accuracy

Error metrics: for some data set $D = \{(t^{(k)}, x^{(k)}), (V^{(k)}, p^{(k)})\}_{k=1}^K$

- $\text{RMAE}_V(\theta; D) := \frac{\sum_{k=1}^K |V^{\text{NN}}(t^{(k)}, x^{(k)}; \theta) - V^{(k)}|}{\sum_{k=1}^K |V^{(k)}|}$
- $\text{RML}^1_p(\theta; D) := \frac{\sum_{k=1}^K \|V_x^{\text{NN}}(t^{(k)}, x^{(k)}; \theta) - p^{(k)}\|_1}{\sum_{k=1}^K \|p^{(k)}\|_1}$
- $\text{RMAE}_u(\theta; D) := \frac{\sum_{k=1}^K |u^{\text{NN}}(t^{(k)}, x^{(k)}; \theta) - u^{(k)}|}{\sum_{k=1}^K |u^{(k)}|}, \quad u^{(k)} = \mathbf{u}(t^{(k)})$

→ performance measure E applied to $D_{\text{train}}, D_{\text{val}}, D_{\text{test}}$

Empirical Validation of Model Accuracy

Error metrics: for some data set $D = \{(t^{(k)}, x^{(k)}), (V^{(k)}, p^{(k)})\}_{k=1}^K$

- $\text{RMAE}_V(\theta; D) := \frac{\sum_{k=1}^K |V^{\text{NN}}(t^{(k)}, x^{(k)}; \theta) - V^{(k)}|}{\sum_{k=1}^K |V^{(k)}|}$
- $\text{RML}^1_p(\theta; D) := \frac{\sum_{k=1}^K \|V_x^{\text{NN}}(t^{(k)}, x^{(k)}; \theta) - p^{(k)}\|_1}{\sum_{k=1}^K \|p^{(k)}\|_1}$
- $\text{RMAE}_u(\theta; D) := \frac{\sum_{k=1}^K |u^{\text{NN}}(t^{(k)}, x^{(k)}; \theta) - u^{(k)}|}{\sum_{k=1}^K |u^{(k)}|}, \quad u^{(k)} = \mathbf{u}(t^{(k)})$

→ performance measure E applied to $D_{\text{train}}, D_{\text{val}}, D_{\text{test}}$

Training and validation error: at epoch i in training round j

$$E_{\text{train}}(i) := E(\theta^i; D_{\text{train}}^j) \quad \text{and} \quad E_{\text{val}}(i) := E(\theta^i; D_{\text{val}})$$

Empirical Validation of Model Accuracy

Error metrics: for some data set $D = \{(t^{(k)}, x^{(k)}), (V^{(k)}, p^{(k)})\}_{k=1}^K$

- $\text{RMAE}_V(\theta; D) := \frac{\sum_{k=1}^K |V^{\text{NN}}(t^{(k)}, x^{(k)}; \theta) - V^{(k)}|}{\sum_{k=1}^K |V^{(k)}|}$
- $\text{RML}^1_p(\theta; D) := \frac{\sum_{k=1}^K \|V_x^{\text{NN}}(t^{(k)}, x^{(k)}; \theta) - p^{(k)}\|_1}{\sum_{k=1}^K \|p^{(k)}\|_1}$
- $\text{RMAE}_u(\theta; D) := \frac{\sum_{k=1}^K |u^{\text{NN}}(t^{(k)}, x^{(k)}; \theta) - u^{(k)}|}{\sum_{k=1}^K |u^{(k)}|}, \quad u^{(k)} = \mathbf{u}(t^{(k)})$

→ performance measure E applied to $D_{\text{train}}, D_{\text{val}}, D_{\text{test}}$

Training and validation error: at epoch i in training round j

$$E_{\text{train}}(i) := E(\theta^i; D_{\text{train}}^j) \quad \text{and} \quad E_{\text{val}}(i) := E(\theta^i; D_{\text{val}})$$

Test error: for final model parameter θ^*

$$E_{\text{test}} := E(\theta^*; D_{\text{test}})$$

Variance Estimation based Sample Size Selection

Goal: find NN parameter $\theta^* \in \mathbb{R}^R$ minimizing population cost function $\text{Cost}(\cdot; D_\infty)$ where $D_\infty := [t_0, t_f] \times \mathbb{X}$

Variance Estimation based Sample Size Selection

- Goal:** find NN parameter $\theta^* \in \mathbb{R}^R$ minimizing population cost function $\text{Cost}(\cdot; D_\infty)$ where $D_\infty := [t_0, t_f] \times \mathbb{X}$
- But:** only sample $D_{\text{train}}^j \subset D_\infty$ of size $K_j := |D_{\text{train}}^j|$ available

Variance Estimation based Sample Size Selection

Goal: find NN parameter $\theta^* \in \mathbb{R}^R$ minimizing population cost function $\text{Cost}(\cdot; D_\infty)$ where $D_\infty := [t_0, t_f] \times \mathbb{X}$

But: only sample $D_{\text{train}}^j \subset D_\infty$ of size $K_j := |D_{\text{train}}^j|$ available

Assumption

All data points $(t^{(k)}, x^{(k)}) \in D_{\text{train}}^j$ are independent and identically distributed.

Variance Estimation based Sample Size Selection

Goal: find NN parameter $\theta^* \in \mathbb{R}^R$ minimizing population cost function $\text{Cost}(\cdot; D_\infty)$ where $D_\infty := [t_0, t_f] \times \mathbb{X}$

But: only sample $D_{\text{train}}^j \subset D_\infty$ of size $K_j := |D_{\text{train}}^j|$ available

Assumption

All data points $(t^{(k)}, x^{(k)}) \in D_{\text{train}}^j$ are independent and identically distributed.

$\rightarrow \nabla \text{Cost}(\theta; D_{\text{train}}^j)$ is unbiased estimator of $\nabla \text{Cost}(\theta; D_\infty)$

Variance Estimation based Sample Size Selection

Goal: find NN parameter $\theta^* \in \mathbb{R}^R$ minimizing population cost function $\text{Cost}(\cdot; D_\infty)$ where $D_\infty := [t_0, t_f] \times \mathbb{X}$

But: only sample $D_{\text{train}}^j \subset D_\infty$ of size $K_j := |D_{\text{train}}^j|$ available

Assumption

All data points $(t^{(k)}, x^{(k)}) \in D_{\text{train}}^j$ are independent and identically distributed.

$\rightarrow \nabla \text{Cost}(\theta; D_{\text{train}}^j)$ is unbiased estimator of $\nabla \text{Cost}(\theta; D_\infty)$

Approach inspired by [Byrd et al., 2012]: control root mean squared error

$$\sqrt{\text{MSE} \left(\nabla \text{Cost}(\theta; D_{\text{train}}^j) \right)} \leq C \left\| \mathbb{E}_{D_{\text{train}}^j} \left[\nabla \text{Cost}(\theta; D_{\text{train}}^j) \right] \right\|_1, \quad C > 0$$

Variance Estimation based Sample Size Selection

Goal: find NN parameter $\theta^* \in \mathbb{R}^R$ minimizing population cost function $\text{Cost}(\cdot; D_\infty)$ where $D_\infty := [t_0, t_f] \times \mathbb{X}$

But: only sample $D_{\text{train}}^j \subset D_\infty$ of size $K_j := |D_{\text{train}}^j|$ available

Assumption

All data points $(t^{(k)}, x^{(k)}) \in D_{\text{train}}^j$ are independent and identically distributed.

$\rightarrow \nabla \text{Cost}(\theta; D_{\text{train}}^j)$ is unbiased estimator of $\nabla \text{Cost}(\theta; D_\infty)$

Approach inspired by [Byrd et al., 2012]: control root mean squared error

$$\sqrt{\frac{1}{K_j} \|\text{Var}_{(t,x) \in D_\infty} (\nabla \text{Cost}(\theta; (t, x)))\|_1} \leq C \left\| \mathbb{E}_{D_{\text{train}}^j} [\nabla \text{Cost}(\theta; D_{\text{train}}^j)] \right\|_1$$

bias-variance decomposition and assumption of i. i. d. data

Variance Estimation based Sample Size Selection

Goal: find NN parameter $\theta^* \in \mathbb{R}^R$ minimizing population cost function $\text{Cost}(\cdot; D_\infty)$ where $D_\infty := [t_0, t_f] \times \mathbb{X}$

But: only sample $D_{\text{train}}^j \subset D_\infty$ of size $K_j := |D_{\text{train}}^j|$ available

Assumption

All data points $(t^{(k)}, x^{(k)}) \in D_{\text{train}}^j$ are independent and identically distributed.

$\rightarrow \nabla \text{Cost}(\theta; D_{\text{train}}^j)$ is unbiased estimator of $\nabla \text{Cost}(\theta; D_\infty)$

Approach inspired by [Byrd et al., 2012]: control root mean squared error

$$\sqrt{\frac{1}{K_j} \left\| \text{Var}_{(t,x) \in D_{\text{train}}^j} (\nabla \text{Cost}(\theta; (t, x))) \right\|_1} \leq C \left\| \nabla \text{Cost}(\theta; D_{\text{train}}^j) \right\|_1$$

approximation by sample variance and sample gradient

Variance Estimation based Sample Size Selection

Convergence test: with tolerance $C > 0$

$$\frac{\sqrt{\left\| \text{Var}_{(t,x) \in D_{\text{train}}^j} \left(\nabla \text{Cost}(\theta; (t, x)) \right) \right\|_1}}{\sqrt{K_j} \left\| \nabla \text{Cost}(\theta; D_{\text{train}}^j) \right\|_1} \leq C$$

Variance Estimation based Sample Size Selection

Convergence test: with tolerance $C > 0$

$$\frac{\sqrt{\left\| \text{Var}_{(t,x) \in D_{\text{train}}^j} \left(\nabla \text{Cost}(\theta; (t, x)) \right) \right\|_1}}{\sqrt{K_j} \left\| \nabla \text{Cost}(\theta; D_{\text{train}}^j) \right\|_1} \leq C$$

Sample size selection rule:

$$K_{j+1} \geq \frac{\left\| \text{Var}_{(t,x) \in D_{\text{train}}^j} \left(\nabla \text{Cost}(\theta; (t, x)) \right) \right\|_1}{C^2 \left\| \nabla \text{Cost}(\theta; D_{\text{train}}^j) \right\|_1^2}$$

Variance Estimation based Sample Size Selection

Convergence test: with tolerance $C > 0$

$$\frac{\sqrt{\left\| \text{Var}_{(t,x) \in D_{\text{train}}^j} \left(\nabla \text{Cost}(\theta; (t, x)) \right) \right\|_1}}{\sqrt{K_j} \left\| \nabla \text{Cost}(\theta; D_{\text{train}}^j) \right\|_1} \leq C$$

Sample size selection rule: with upper bound $M > 1$

$$K_{j+1} := \min \left\{ \left\lceil \frac{\left\| \text{Var}_{(t,x) \in D_{\text{train}}^j} \left(\nabla \text{Cost}(\theta; (t, x)) \right) \right\|_1}{C^2 \left\| \nabla \text{Cost}(\theta; D_{\text{train}}^j) \right\|_1^2} \right\rceil, \lfloor MK_j \rfloor \right\}$$

Early Stopping

Preventing excessive computation of convergence test and data generation

- delay until overfitting is detected by early stopping criterion
- promote cheaper heuristics and efficient use of data
- reduce computational costs during training phase

Early Stopping

Preventing excessive computation of convergence test and data generation

- delay until overfitting is detected by early stopping criterion
- promote cheaper heuristics and efficient use of data
- reduce computational costs during training phase

Criterion adapted from [Prechelt, 2012]: with generalization loss GL, training progress TP_ι over ι epochs, and threshold $T > 0$

$$\text{stop at epoch } i \geq \iota \text{ if: } TP_\iota(i) = 0 \quad \text{or} \quad \frac{GL(i)}{TP_\iota(i)} > T$$

Algorithm Adaptive Deep Learning

1. Generate D_{train}^0 and D_{val}
 2. Initialize NN with $\theta \in \mathbb{R}^R$; set $i \leftarrow 0$, $j \leftarrow 0$, $i^* \leftarrow 0$ and $\theta^* \leftarrow \theta$
 - while** not converged **do**
 - while** not early stopped **do**
 3. Update θ by training NN on D_{train}^j for one epoch; set $i \leftarrow i + 1$
 4. Track $E_{\text{train}}(i)$ and $E_{\text{val}}(i)$; evaluate early stopping criterion
 - if** $E_{\text{val}}(i) < E_{\text{val}}(i^*)$ **then**
 5. Set $i^* \leftarrow i$ and $\theta^* \leftarrow \theta$
 6. Evaluate convergence criterion
 - if** not converged **then**
 7. Determine $|D_{\text{train}}^{j+1}|$ by sample size criterion
 8. Augment D_{train}^j adaptively; set $j \leftarrow j + 1$
 9. Load NN with θ^*
 10. Generate D_{test}
 11. Compute E_{test} to evaluate generalization performance
-

Algorithm Adaptive Deep Learning

1. Generate D_{train}^0 and D_{val}
 2. Initialize NN with $\theta \in \mathbb{R}^R$; set $i \leftarrow 0$, $j \leftarrow 0$, $i^* \leftarrow 0$ and $\theta^* \leftarrow \theta$
 - while** not converged **do**
 - while** not early stopped **do**
 3. Update θ by training NN on D_{train}^j for one epoch; set $i \leftarrow i + 1$
 4. Track $E_{\text{train}}(i)$ and $E_{\text{val}}(i)$; evaluate early stopping criterion
 - if** $E_{\text{val}}(i) < E_{\text{val}}(i^*)$ **then**
 5. Set $i^* \leftarrow i$ and $\theta^* \leftarrow \theta$
 6. Evaluate convergence criterion
 - if** not converged **then**
 7. Determine $|D_{\text{train}}^{j+1}|$ by sample size criterion
 8. Augment D_{train}^j adaptively; set $j \leftarrow j + 1$
 9. Load NN with θ^*
 10. Generate D_{test}
 11. Compute E_{test} to evaluate generalization performance
-

Algorithm Adaptive Deep Learning

1. Generate D_{train}^0 and D_{val}
 2. Initialize NN with $\theta \in \mathbb{R}^R$; set $i \leftarrow 0$, $j \leftarrow 0$, $i^* \leftarrow 0$ and $\theta^* \leftarrow \theta$
 - while** not converged **do**
 - while** not early stopped **do**
 3. Update θ by training NN on D_{train}^j for one epoch; set $i \leftarrow i + 1$
 4. Track $E_{\text{train}}(i)$ and $E_{\text{val}}(i)$; evaluate early stopping criterion
 - if** $E_{\text{val}}(i) < E_{\text{val}}(i^*)$ **then**
 5. Set $i^* \leftarrow i$ and $\theta^* \leftarrow \theta$
 6. Evaluate convergence criterion
 - if** not converged **then**
 7. Determine $|D_{\text{train}}^{j+1}|$ by sample size criterion
 8. Augment D_{train}^j adaptively; set $j \leftarrow j + 1$
 9. Load NN with θ^*
 10. Generate D_{test}
 11. Compute E_{test} to evaluate generalization performance
-

Algorithm Adaptive Deep Learning

1. Generate D_{train}^0 and D_{val}
 2. Initialize NN with $\theta \in \mathbb{R}^R$; set $i \leftarrow 0$, $j \leftarrow 0$, $i^* \leftarrow 0$ and $\theta^* \leftarrow \theta$
 - while** not converged **do**
 - while** not early stopped **do**
 3. Update θ by training NN on D_{train}^j for one epoch; set $i \leftarrow i + 1$
 4. Track $E_{\text{train}}(i)$ and $E_{\text{val}}(i)$; evaluate **early stopping criterion**
 - if** $E_{\text{val}}(i) < E_{\text{val}}(i^*)$ **then**
 5. Set $i^* \leftarrow i$ and $\theta^* \leftarrow \theta$
 6. Evaluate **convergence criterion**
 - if** not converged **then**
 7. Determine $|D_{\text{train}}^{j+1}|$ by **sample size criterion**
 8. Augment D_{train}^j adaptively; set $j \leftarrow j + 1$
 9. Load NN with θ^*
 10. Generate D_{test}
 11. Compute E_{test} to evaluate generalization performance
-

Algorithm Adaptive Deep Learning

1. Generate D_{train}^0 and D_{val}
 2. Initialize NN with $\theta \in \mathbb{R}^R$; set $i \leftarrow 0$, $j \leftarrow 0$, $i^* \leftarrow 0$ and $\theta^* \leftarrow \theta$
 - while** not converged **do**
 - while** not early stopped **do**
 3. Update θ by training NN on D_{train}^j for one epoch; set $i \leftarrow i + 1$
 4. Track $E_{\text{train}}(i)$ and $E_{\text{val}}(i)$; evaluate early stopping criterion
 - if** $E_{\text{val}}(i) < E_{\text{val}}(i^*)$ **then**
 5. Set $i^* \leftarrow i$ and $\theta^* \leftarrow \theta$
 6. Evaluate convergence criterion
 - if** not converged **then**
 7. Determine $|D_{\text{train}}^{j+1}|$ by sample size criterion
 8. Augment D_{train}^j adaptively; set $j \leftarrow j + 1$
 9. Load NN with θ^*
 10. Generate D_{test}
 11. Compute E_{test} to evaluate generalization performance
-

Algorithm Adaptive Deep Learning

1. Generate D_{train}^0 and D_{val}
 2. Initialize NN with $\theta \in \mathbb{R}^R$; set $i \leftarrow 0$, $j \leftarrow 0$, $i^* \leftarrow 0$ and $\theta^* \leftarrow \theta$
 - while** not converged **do**
 - while** not early stopped **do**
 3. Update θ by training NN on D_{train}^j for one epoch; set $i \leftarrow i + 1$
 4. Track $E_{\text{train}}(i)$ and $E_{\text{val}}(i)$; evaluate early stopping criterion
 - if** $E_{\text{val}}(i) < E_{\text{val}}(i^*)$ **then**
 5. Set $i^* \leftarrow i$ and $\theta^* \leftarrow \theta$
 6. Evaluate convergence criterion
 - if** not converged **then**
 7. Determine $|D_{\text{train}}^{j+1}|$ by sample size criterion
 8. Augment D_{train}^j adaptively; set $j \leftarrow j + 1$
 9. Load NN with θ^*
 10. Generate D_{test}
 11. Compute E_{test} to evaluate generalization performance
-

- 1 Introduction
- 2 Optimal Control Theory
- 3 Adaptive Deep Learning
- 4 Application to Reaction-Diffusion System**
- 5 Numerical Experiments
- 6 Conclusion and Future Research

Optimal Control of Reaction-Diffusion System

Diffusion equation with state $X : [t_0, t_f] \times \Omega \rightarrow \mathbb{R}$ on $\Omega \subseteq \mathbb{R}$, $a > 0$

$$X_t(t, \omega) = a \Delta X(t, \omega) \quad \text{for a. a. } (t, \omega) \in [t_0, t_f] \times \Omega$$

Optimal Control of Reaction-Diffusion System

Reaction term with $b > 0$

$$X_t(t, \omega) = a\Delta X(t, \omega) + bR(X(t, \omega)) \quad \text{for a. a. } (t, \omega)$$

introducing non-linearity via, e. g.,

- hyperbolic growth rate $R_1(X(t, \omega)) := X(t, \omega)^2$
- logistic growth rate $R_2(X(t, \omega)) := X(t, \omega)(1 - X(t, \omega))$

Optimal Control of Reaction-Diffusion System

Control term with control $\mathbf{u} : [t_0, t_f] \rightarrow \mathbb{R}$, $c > 0$

$$X_t(t, \omega) = a\Delta X(t, \omega) + bR(X(t, \omega)) + cS(\omega, \mathbf{u}(t), X(t, \omega)) \quad \text{for a. a. } (t, \omega)$$

as source or sink via, e. g.,

- additive control $S_1(\omega, \mathbf{u}(t)) := \mathbf{1}_{\Omega_S}(\omega)\mathbf{u}(t)$, $\Omega_S \subseteq \Omega$
- bilinear control $S_2(\mathbf{u}(t), X(t, \omega)) := \mathbf{u}(t)X(t, \omega)$

Optimal Control of Reaction-Diffusion System

Controlled reaction-diffusion system

$$X_t(t, \omega) = a\Delta X(t, \omega) + bR(X(t, \omega)) + cS(\omega, \mathbf{u}(t), X(t, \omega)) \quad \text{for a. a. } (t, \omega)$$

Dirichlet boundary conditions with initial condition $X_0 : \Omega \rightarrow \mathbb{R}$

$$\begin{cases} X(t, \omega) = 0 & \text{on } [t_0, t_f] \times \partial\Omega \\ X(t_0, \omega) = X_0(\omega) & \text{on } \Omega, \end{cases}$$

Optimal Control of Reaction-Diffusion System

Controlled reaction-diffusion system

$$X_t(t, \omega) = a\Delta X(t, \omega) + bR(X(t, \omega)) + cS(\omega, \mathbf{u}(t), X(t, \omega)) \quad \text{for a. a. } (t, \omega)$$

Dirichlet boundary conditions

$$\begin{cases} X(t, \omega) = 0 & \text{on } [t_0, t_f] \times \partial\Omega \\ X(t_0, \omega) = X_0(\omega) & \text{on } \Omega, \end{cases}$$

Quadratic cost function with $\alpha, \beta, \gamma > 0$

$$J(\mathbf{u}, X) = \int_{t_0}^{t_f} \left[\alpha \|X(t, \cdot) - Y(\cdot)\|_{L^2(\Omega)}^2 + \beta |\mathbf{u}(t)|^2 \right] dt \\ + \gamma \|X(t_f, \cdot) - Y(\cdot)\|_{L^2(\Omega)}^2$$

steering X to target $Y : \Omega \rightarrow \mathbb{R}$ expending minimum control effort

Optimal Control of Reaction-Diffusion System

Controlled reaction-diffusion system

$$X_t(t, \omega) = a\Delta X(t, \omega) + bR(X(t, \omega)) + cS(\omega, \mathbf{u}(t), X(t, \omega)) \quad \text{for a. a. } (t, \omega)$$

Dirichlet boundary conditions

$$\begin{cases} X(t, \omega) = 0 & \text{on } [t_0, t_f] \times \partial\Omega \\ X(t_0, \omega) = X_0(\omega) & \text{on } \Omega, \end{cases}$$

Quadratic cost function

$$J(\mathbf{u}, X) = \int_{t_0}^{t_f} \left[\alpha \|X(t, \cdot) - Y(\cdot)\|_{L^2(\Omega)}^2 + \beta |\mathbf{u}(t)|^2 \right] dt \\ + \gamma \|X(t_f, \cdot) - Y(\cdot)\|_{L^2(\Omega)}^2$$

→ transforming PDE- into ODE-constrained OCP via spatial discretization

- 1 Introduction
- 2 Optimal Control Theory
- 3 Adaptive Deep Learning
- 4 Application to Reaction-Diffusion System
- 5 Numerical Experiments**
- 6 Conclusion and Future Research

Implementation and Experimental Set-up

Modular code for discretized OCP of reaction-diffusion system:

PROBLEM

time interval with $t_0 = 0$ and $t_f = 5$,
domain $\Omega = [0, 1]$ discretized with $n = 20$ mesh nodes,
reaction-diffusion system and quadratic cost function with respective coefficients

DATA GENERATION

initial condition domain
 $\mathbb{X}_0 = [-1.5, 1.5]^n$,
upper bound $M = 1.25$

MODEL

ResNet of depth $L = 5$ and width $d = 100$

TRAINING

Adam with gradient regularization $\lambda = 100$,
convergence tolerance $C = 10^{-4}$

EVALUATION

monitoring of training and test errors, data set size and runtime,
simulations of NN and LQR controllers in presence of noise

Numerical experiments:


- ① Noise in the Linear-Quadratic Problem
- ② Non-linearities in the System Dynamics
- ③ Heuristic-Guided Training for Bilinear Control
- ④ Scalability across System Dimensions
- ⑤ BVP Solver Initialization for Data Generation Efficiency
- ⑥ Simulation for Out-of-Domain Initial Conditions

Numerical experiments:

- 1 Noise in the Linear-Quadratic Problem
- 2 Non-linearities in the System Dynamics
- 3 Heuristic-Guided Training for Bilinear Control
- 4 Scalability across System Dimensions
- 5 BVP Solver Initialization for Data Generation Efficiency
- 6 Simulation for Out-of-Domain Initial Conditions

Noise in the Linear-Quadratic Problem


- consider linear-quadratic case, i. e., additive control without reaction
- analyse effect of Gaussian and shock noise
- simulate NN and LQR feedback, as well as open-loop BVP solution



solution technique	total cost
NN	1.56
LQR	1.65
BVP	6.21

Noise in the Linear-Quadratic Problem

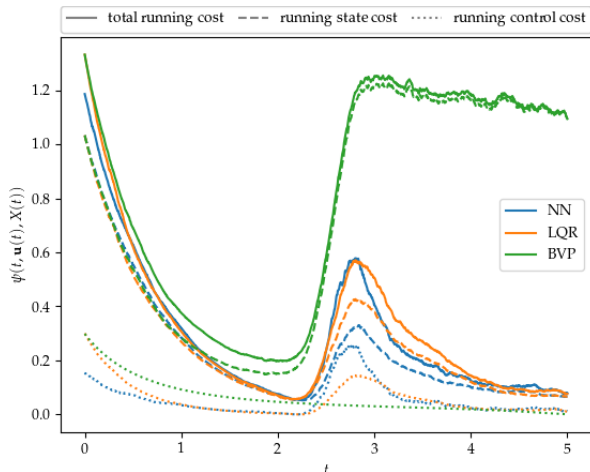
→ open-loop control is unable to respond to disturbances



solution technique	total cost
NN	1.56
LQR	1.65
BVP	6.21

Noise in the Linear-Quadratic Problem

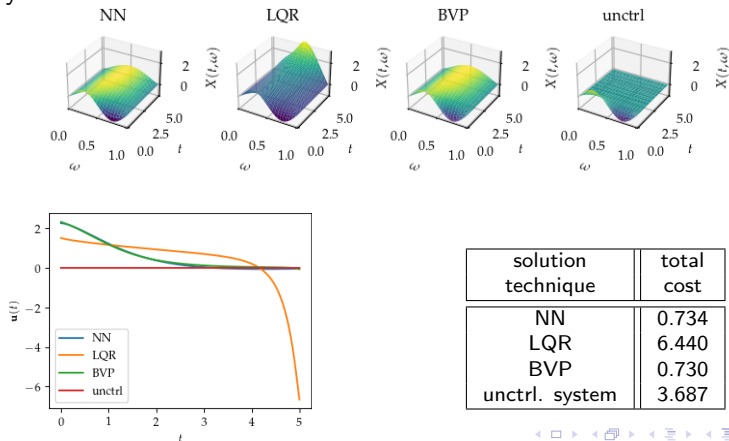
- open-loop control is unable to respond to disturbances
- NN and LQR feedback are robust to stochastic and deterministic noise, NN even more than LQR



solution technique	total cost
NN	1.56
LQR	1.65
BVP	6.21

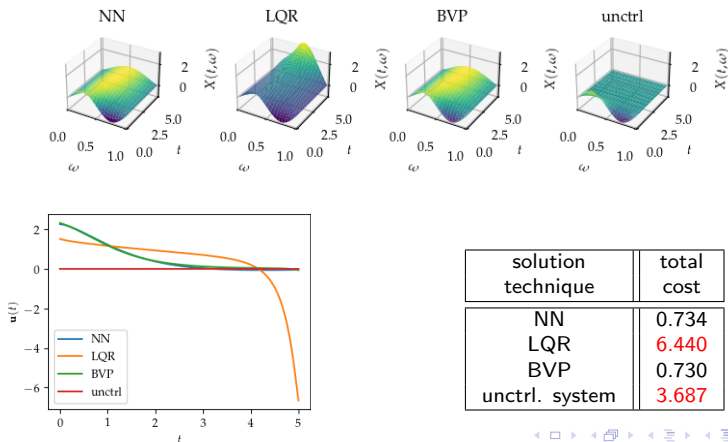
Non-linearities in the System Dynamics

- add non-linear reaction term to system, e. g., hyperbolic growth rate
- apply LQR to linearized state dynamics
- benchmark against optimal open-loop solution and uncontrolled system



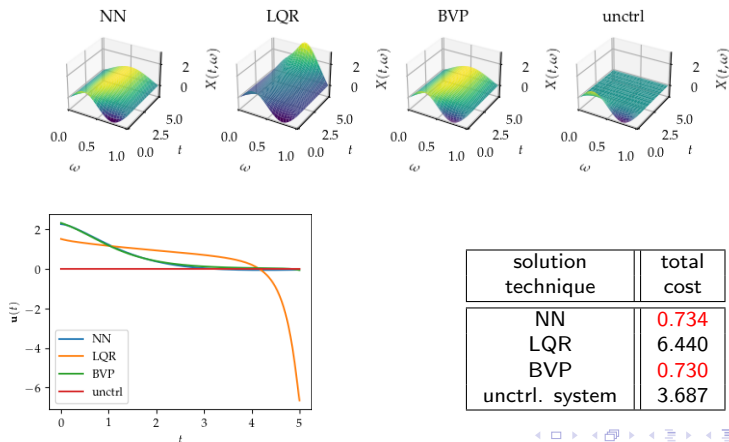
Non-linearities in the System Dynamics

→ linearization based LQR may even perform worse than uncontrolled system




Non-linearities in the System Dynamics

- linearization based LQR may even perform worse than uncontrolled system
- NN solution closely resembles optimal solution provided by BVP



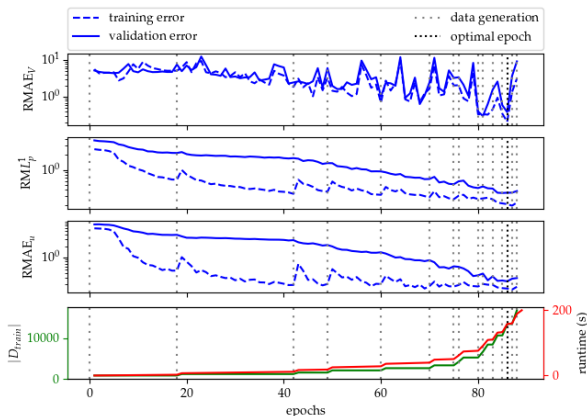
Heuristic-Guided Training for Bilinear Control

- replace additive by bilinear control
- monitor errors, amount of data and runtime during training phase
- evaluate effects of adaptive and progressive data generation




Heuristic-Guided Training for Bilinear Control

- data expansion is delayed until significant increase of generalization error and becomes more frequent for model refinement



Heuristic-Guided Training for Bilinear Control

- data expansion is delayed until significant increase of generalization error and becomes more frequent for model refinement
- combination of stochastic optimization, gradient regularization, early stopping, dynamic data set size and adaptive sampling promises effective training



Scalability across System Dimensions

- refine spatial discretization yielding high dimensional states
- compute test errors for empirical verification of NN accuracy

limited training time of 100 s
(hyperbolic bilinear problem setting
with zero target)

n	# data points	RMAE_u
5	11306	0.24
20	6626	0.27
40	4895	2.39

satisfied convergence criterion
(hyperbolic bilinear problem setting
with quadratic target)

n	training time	RMAE_u
5	113 s	0.93
20	159 s	0.87
40	502 s	1.17

Scalability across System Dimensions

- refine spatial discretization yielding high dimensional states
- compute test errors for empirical verification of NN accuracy

limited training time of 100 s
(hyperbolic bilinear problem setting
with zero target)

n	# data points	RMAE_u
5	11306	0.24
20	6626	0.27
40	4895	2.39

satisfied convergence criterion
(hyperbolic bilinear problem setting
with quadratic target)

n	training time	RMAE_u
5	113 s	0.93
20	159 s	0.87
40	502 s	1.17

→ accuracy declines as
dimension grows

Scalability across System Dimensions

- refine spatial discretization yielding high dimensional states
- compute test errors for empirical verification of NN accuracy

limited training time of 100 s
(hyperbolic bilinear problem setting
with zero target)

n	# data points	RMAE_u
5	11306	0.24
20	6626	0.27
40	4895	2.39

satisfied convergence criterion
(hyperbolic bilinear problem setting
with quadratic target)

n	training time	RMAE_u
5	113 s	0.93
20	159 s	0.87
40	502 s	1.17

- accuracy declines as dimension grows
- less data-intensive due to slower overfitting

Scalability across System Dimensions

- refine spatial discretization yielding high dimensional states
- compute test errors for empirical verification of NN accuracy

limited training time of 100 s
(hyperbolic bilinear problem setting
with zero target)

n	# data points	RMAE_u
5	11306	0.24
20	6626	0.27
40	4895	2.39

- accuracy declines as dimension grows
- less data-intensive due to slower overfitting

satisfied convergence criterion
(hyperbolic bilinear problem setting
with quadratic target)

n	training time	RMAE_u
5	113 s	0.93
20	159 s	0.87
40	502 s	1.17

- high-fidelity models across all dimensions

Scalability across System Dimensions

- refine spatial discretization yielding high dimensional states
- compute test errors for empirical verification of NN accuracy

limited training time of 100 s
(hyperbolic bilinear problem setting
with zero target)

n	# data points	RMAE_u
5	11306	0.24
20	6626	0.27
40	4895	2.39

- accuracy declines as dimension grows
- less data-intensive due to slower overfitting

satisfied convergence criterion
(hyperbolic bilinear problem setting
with quadratic target)

n	training time	RMAE_u
5	113 s	0.93
20	159 s	0.87
40	502 s	1.17

- high-fidelity models across all dimensions
- computational effort scales reasonably

Scalability across System Dimensions

- refine spatial discretization yielding high dimensional states
- compute test errors for empirical verification of NN accuracy

limited training time of 100 s
(hyperbolic bilinear problem setting
with zero target)

n	# data points	RMAE_u
5	11306	0.24
20	6626	0.27
40	4895	2.39

satisfied convergence criterion
(hyperbolic bilinear problem setting
with quadratic target)

n	training time	RMAE_u
5	113 s	0.93
20	159 s	0.87
40	502 s	1.17

→ DL method remains viable at least up to problem dimension 40

Outline

- 1 Introduction
- 2 Optimal Control Theory
- 3 Adaptive Deep Learning
- 4 Application to Reaction-Diffusion System
- 5 Numerical Experiments
- 6 Conclusion and Future Research**

Conclusion

- address challenge of feedback design for high-dimensional non-linear optimal control problems
- establish adaptive deep learning framework for solving Hamilton-Jacobi-Bellman equations
- conduct numerical analysis demonstrating scalability and robustness of neural network controllers

Conclusion and Future Research

Conclusion

- address challenge of feedback design for high-dimensional non-linear optimal control problems
- establish adaptive deep learning framework for solving Hamilton-Jacobi-Bellman equations
- conduct numerical analysis demonstrating scalability and robustness of neural network controllers

Future Research

- develop theoretical results to further enhance performance of neural network feedback
- investigate alternative models, optimization techniques and data generation methods
- refine heuristics for progressive data augmentation dependent on specific application
- tackle infinite-horizon or minimum-time problems, with state and control constraints, or infinite dimensional systems

References I

 Byrd, R., Chin, G., Nocedal, J., and Wu, Y. (2012).

Sample Size Selection in Optimization Methods for Machine Learning.
Mathematical Programming, 134:127–155.

 Kang, W. and Wilcox, L. C. (2015).

Mitigating the Curse of Dimensionality: Sparse Grid Characteristics Method for Optimal Feedback Control and HJB Equations.

 Kunisch, Karl and Walter, Daniel (2021).

Semiglobal optimal feedback stabilization of autonomous systems via deep neural network approximation.
ESAIM: COCV, 27:16.

 Liberzon, D. (2012).

Calculus of Variations and Optimal Control Theory: A Concise Introduction.
Princeton University Press.

References II



Nakamura-Zimmerer, T., Gong, Q., and Kang, W. (2021).
Adaptive Deep Learning for High-Dimensional Hamilton–Jacobi–Bellman
Equations.
SIAM Journal on Scientific Computing, 43(2):A1221–A1247.



Prechelt, L. (2012).
Early Stopping — But When?
Springer Berlin Heidelberg.



Raissi, M., Perdikaris, P., and Karniadakis, G. E. (2019).
Physics-informed neural networks: A deep learning framework for solving forward
and inverse problems involving nonlinear partial differential equations.
Journal of Computational Physics, 378:686–707.



Schättler, H. and Ledzewicz, U. (2012).
Geometric Optimal Control: Theory, Methods and Examples.
Interdisciplinary Applied Mathematics. Springer New York.

Thank you for your attention!

I am now available to answer any questions you may have.

Early Stopping

- performance measure E
for tracking training and validation error at epoch i in training round j

$$E_{\text{train}}(i) := E(\theta^i; D_{\text{train}}^j) \quad \text{and} \quad E_{\text{val}}(i) := E(\theta^i; D_{\text{val}}),$$

determining generalization loss

$$\text{GL}(i) := \frac{E_{\text{val}}(i)}{\min_{1 \leq i' \leq i} E_{\text{val}}(i')} - 1$$

and training progress for training strip of length ι

$$\text{TP}_{\iota}(i) := \frac{\sum_{i'=i-\iota+1}^i E_{\text{train}}(i')}{\iota \cdot \min_{i-\iota+1 \leq i' \leq i} E_{\text{train}}(i')} - 1$$

- trigger with threshold $T > 0$
for detecting overfitting and interrupting training

$$\text{stop at epoch } i \geq \iota \text{ if: } \text{TP}_{\iota}(i) = 0 \quad \text{or} \quad \frac{\text{GL}(i)}{\text{TP}_{\iota}(i)} > T$$

Key Strengths of AdaDL Algorithm

- exploiting prior information:
value function BC, NN warm start, physics-informed learning
- using data efficiently:
costate data usage, adaptive sampling, data generation delay
- promoting generalization:
early stopping, gradient and L2 regularization, stochastic optimization
- reducing the sensitivity to hyperparameter tuning:
Adam, ResNet, data set size heuristic

Spatial Discretization

- using central difference scheme on equidistant mesh $\{\omega_j\}_{j=1}^n \subset \Omega$ of size h

$$\dot{\mathbf{x}}_j(t) = \frac{a}{h^2}(\mathbf{x}_{j-1}(t) - 2\mathbf{x}_j(t) + \mathbf{x}_{j+1}(t)) + bR(\mathbf{x}_j(t)) + cS(\omega_j, \mathbf{u}(t), \mathbf{x}_j(t))$$

for a. a. t and all $j = 1, \dots, n$

→ dimension n of state $\mathbf{x} := (\mathbf{x}_1, \dots, \mathbf{x}_n)^\top$ is scalable

- transforming into ODE-constrained OCP

$$\min_{(\mathbf{u}, \mathbf{x}) \in \mathcal{U} \times \mathcal{X}} \int_{t_0}^{t_f} \left[\alpha h \|\mathbf{x}(t) - y\|_2^2 + \beta |\mathbf{u}(t)|^2 \right] dt + \gamma h \|\mathbf{x}(t_f) - y\|_2^2$$

subject to $\dot{\mathbf{x}}(t) = aQ\mathbf{x}(t) + bR(\mathbf{x}(t)) + cS(\mathbf{u}(t), \mathbf{x}(t))$ for a. a. t ;
 $\mathbf{x}(t_0) = x_0$

Boundary Value Problem and Neural Network Control

Boundary value problem specified for discretized OCP

$$\left\{ \begin{array}{l} \dot{\mathbf{x}}(t) = aQ\mathbf{x}(t) + bR(\mathbf{x}(t)) + cS(\mathbf{u}(t), \mathbf{x}(t)) \\ \quad \text{for a. a. } t; \quad \mathbf{x}(t_0) = x_0, \\ \dot{\mathbf{p}}(t) = -2\alpha h(\mathbf{x}(t) - y) - aQ^\top \mathbf{p}(t) - bR_x^\top(\mathbf{x}(t))\mathbf{p}(t) - cS_x^\top(\mathbf{x}(t))\mathbf{p}(t) \\ \quad \text{for a. a. } t; \quad \mathbf{p}(t_f) = 2\gamma h(\mathbf{x}(t_f) - y), \\ \dot{\mathbf{v}}(t) = -\alpha h \|\mathbf{x}(t) - y\|_2^2 - \beta |\mathbf{u}(t)|^2 \\ \quad \text{for a. a. } t; \quad \mathbf{v}(t_f) = \gamma h \|\mathbf{x}(t_f) - y\|_2^2 \end{array} \right.$$

Boundary Value Problem and Neural Network Control

Boundary value problem specified for discretized OCP

$$\left\{ \begin{array}{l} \dot{\mathbf{x}}(t) = aQ\mathbf{x}(t) + bR(\mathbf{x}(t)) + cS(\mathbf{u}(t), \mathbf{x}(t)) \\ \quad \text{for a. a. } t; \quad \mathbf{x}(t_0) = x_0, \\ \dot{\mathbf{p}}(t) = -2\alpha h(\mathbf{x}(t) - y) - aQ^\top \mathbf{p}(t) - bR_x^\top(\mathbf{x}(t))\mathbf{p}(t) - cS_x^\top(\mathbf{x}(t))\mathbf{p}(t) \\ \quad \text{for a. a. } t; \quad \mathbf{p}(t_f) = 2\gamma h(\mathbf{x}(t_f) - y), \\ \dot{\mathbf{v}}(t) = -\alpha h \|\mathbf{x}(t) - y\|_2^2 - \beta |\mathbf{u}(t)|^2 \\ \quad \text{for a. a. } t; \quad \mathbf{v}(t_f) = \gamma h \|\mathbf{x}(t_f) - y\|_2^2 \end{array} \right.$$

Candidate open-loop solution by Hamiltonian minimization condition

$$\mathbf{u}^*(t; x_0) \in \underset{u \in \mathbb{R}}{\operatorname{argmin}} \{ \beta u^2 + c\mathbf{p}^*(t) \cdot S(u, \mathbf{x}^*(t)) \} \quad \text{for a. a. } t$$

Boundary Value Problem and Neural Network Control

Boundary value problem specified for discretized OCP

$$\left\{ \begin{array}{l} \dot{\mathbf{x}}(t) = aQ\mathbf{x}(t) + bR(\mathbf{x}(t)) + cS(\mathbf{u}(t), \mathbf{x}(t)) \\ \quad \text{for a. a. } t; \quad \mathbf{x}(t_0) = x_0, \\ \dot{\mathbf{p}}(t) = -2\alpha h(\mathbf{x}(t) - y) - aQ^\top \mathbf{p}(t) - bR_x^\top(\mathbf{x}(t))\mathbf{p}(t) - cS_x^\top(\mathbf{x}(t))\mathbf{p}(t) \\ \quad \text{for a. a. } t; \quad \mathbf{p}(t_f) = 2\gamma h(\mathbf{x}(t_f) - y), \\ \dot{\mathbf{v}}(t) = -\alpha h \|\mathbf{x}(t) - y\|_2^2 - \beta |\mathbf{u}(t)|^2 \\ \quad \text{for a. a. } t; \quad \mathbf{v}(t_f) = \gamma h \|\mathbf{x}(t_f) - y\|_2^2 \end{array} \right.$$

Neural Network feedback controller

$$\mathbf{u}^{\text{NN}}(t, x) \in \underset{u \in \mathbb{R}}{\operatorname{argmin}} \{ \beta u^2 + cV_x^{\text{NN}}(t, x) \cdot S(u, x) \} \quad \text{for all } t \text{ and } x$$

Linear-Quadratic Regulator

LQR feedback controller computed via Riccati differential equation
→ as baseline for evaluating NN

Linear-Quadratic Regulator

LQR feedback controller computed via Riccati differential equation
→ as baseline for evaluating NN

- consider $R \equiv 0$ and $S = S_1$, i. e., linear system

$$\dot{\mathbf{x}}(t) = a_Q \mathbf{x}(t) + c \mathbb{1}_{\Omega_S} \mathbf{u}(t)$$

→ analytical solution for linear-quadratic problem

Linear-Quadratic Regulator

LQR feedback controller computed via Riccati differential equation
→ as baseline for evaluating NN

- consider $R \equiv 0$ and $S = S_1$, i. e., linear system

$$\dot{\mathbf{x}}(t) = aQ\mathbf{x}(t) + c\mathbb{1}_{\Omega_S}\mathbf{u}(t)$$

→ analytical solution for linear-quadratic problem

- consider **non-linear** $R = R_1$ or R_2 and $S = S_1$, i. e.,

$$\dot{\mathbf{x}}(t) = aQ\mathbf{x}(t) + \mathbf{b}\mathbf{x}(t) \odot \mathbf{x}(t) + c\mathbb{1}_{\Omega_S}\mathbf{u}(t) \quad \text{or}$$

$$\dot{\mathbf{x}}(t) = aQ\mathbf{x}(t) + \mathbf{b}\mathbf{x}(t) \odot (\mathbb{1} - \mathbf{x}(t)) + c\mathbb{1}_{\Omega_S}\mathbf{u}(t)$$

Linear-Quadratic Regulator

LQR feedback controller computed via Riccati differential equation
→ as baseline for evaluating NN

- consider $R \equiv 0$ and $S = S_1$, i. e., linear system

$$\dot{\mathbf{x}}(t) = aQ\mathbf{x}(t) + c\mathbb{1}_{\Omega_S}\mathbf{u}(t)$$

→ analytical solution for linear-quadratic problem

- consider **linearized** $R = R_1$ or R_2 and $S = S_1$, i. e.,

$$\dot{\mathbf{x}}(t) = aQ\mathbf{x}(t) + c\mathbb{1}_{\Omega_S}\mathbf{u}(t) \quad \text{or}$$

$$\dot{\mathbf{x}}(t) = (aQ + bI)\mathbf{x}(t) + c\mathbb{1}_{\Omega_S}\mathbf{u}(t)$$


→ approximate solution via linearization of state dynamics

BVP Solver Initialization for Data Generation Efficiency

- investigate dependency of time-marching on time sequence and NN warm start on model accuracy
 - determine initialization strategy, for problems of varying difficulty and at different training stages, yielding best trade-off of convergence rate and speed
- improved convergence attained through sophisticated techniques, i. e., time-marching with tuned adaptive intervals and NN warm start even with briefly trained low-fidelity models
- speed of NN warm start scales much better with problem dimension than time-marching

Simulation for Out-of-Domain Initial Conditions

- confront NN with data outside its training domain by upscaling the initial state value



solution technique	total cost
NN	3.02
BVP	1.66

- NN model produces acceptable feedback controls, preventing blow-ups but with increased costs
- DL method may handle scenarios in which disturbances propel state beyond learned domain